

Priyansh Vaish

priyansh.vaish01@gmail.com | LinkedIn | GitHub | Portfolio

PROFESSIONAL SUMMARY

Computer Science undergraduate specializing in Artificial Intelligence and Machine Learning, with a focus on computer vision and applied deep learning using YOLOv8, OpenCV, and PyTorch. Experienced in building real-time systems end-to-end, from model design to deployment, with hands-on contributions to open-source software through national academic initiatives.

EDUCATION

VIT Bhopal University

B.Tech in Computer Science (AI & ML) | CGPA: 8.67

Bhopal, India

2023 – 2027

TECHNICAL SKILLS

Languages: Python, C++, HTML5, CSS3, SQL

Data & ML: NumPy, Pandas, scikit-learn, RAG, LLM Applications, Neural Networks, TensorFlow, PyTorch

Computer Vision: OpenCV, YOLOv8, MediaPipe, Tesseract OCR, Depth Estimation, Image Processing

Cloud & DevOps: AWS (EC2, S3, RDS), Docker, Docker Compose, FastAPI, Vercel, Supabase, Git

Tools & Frameworks: Claude Code, Github

Relevant Coursework: Data Structures & Algorithms, Computer Networks, Machine Learning, Computer Vision, DBMS

EXPERIENCE

FOSSEE – OSDAG, IIT Bombay

May 2026 – Present

Summer Intern

Mumbai, India

- Contributing as an open-source developer to OSDAG's 3psLCCA platform, a structural engineering design tool built under FOSSEE, IIT Bombay's national initiative for open-source education software.
- Delivering feature improvements across multiple functional modules of the platform, including carbon-emission tracking and construction work data workflows, within an actively maintained, version-controlled codebase.
- Following a structured open-source development workflow – debugging, testing, and integrating new functionality within a development fork ahead of merge into the main repository.

PROJECTS

EchoVision – Real-Time Edge Vision Assistant | [GitHub](#)

Python /Depth-Anything V2 /OpenCV /Piper TTS

- Depth-based hazard detection using Depth-Anything V2; 30 FPS at 640×480, <200 ms latency.
- 3-region spatial engine (L/C/R) with Piper neural TTS and 15/30/45 degree turn guidance.
- Implemented modular Python architecture with cooldown logic to suppress duplicate real-time alerts.

FaceID Attendance System | [GitHub](#)

Python /OpenCV /face_recognition /Supabase

- Real-time face recognition system (OpenCV): 97% accuracy, sub-200 ms per-frame latency.
- Offline encoding pipeline with session-scoped deduplication, eliminating 95% of duplicate entries.
- Dual-storage architecture (local CSV + Supabase PostgreSQL) enabling multi-location attendance tracking.

Automatic Number Plate Recognition (ANPR) | [GitHub](#)

Python /YOLOv8 /Tesseract OCR /OpenCV

- YOLOv8 + Tesseract OCR pipeline: 90%+ accuracy on 2,000+ video frames at 30 FPS.
- 95% consistent plate localization across varied lighting conditions and angles.
- Modular design with decoupled detection and OCR stages for scalable, high-throughput processing.

ACHIEVEMENTS

Finalist – Inya.ai Buildathon (NASSCOM Initiative)

Oct 2025

Team Dikastirio | Track: Airlines Booking and Status

- Built Hermes, a bilingual AI flight assistant with Twilio-integrated voice/SMS support for bookings, cancellations, and live status queries – finalist among teams nationwide.

CERTIFICATIONS & ACTIVITIES

- **NPTEL, IIT Kharagpur:** Cloud Computing (2025); Introduction to Internet of Things (2026)
- **Machine Learning Specialization** (3-Course Program) – DeepLearning.AI & Stanford University
- Design Team Member – AWS Cloud Club, VIT Bhopal University (2024–2025)